
AI-Driven High-Temperature Material Discovery

An End-to-End ML Pipeline for Ceramic Composition Ranking

Hợp kim gốm chịu nhiệt
Phát hiện vật liệu bằng ML / End-to-End

Contest Report — Contest Submission
AI Material Discovery

AI Material Team
ai-material v0.1 | July 2026
Repository: `research/phase_2/`

Keywords: *generative materials science, GraphVAE, CHGNet relaxation, ML surrogate, refractory carbides, ML4Science*

Contents

Executive Summary / Tóm tắt điều hành	5
I Technical Report / Báo cáo kỹ thuật	6
1 Problem statement / Phát biểu bài toán	7
1.1 Motivation / Động lực	7
1.2 Problem definition / Định nghĩa bài toán	7
1.3 Scientific limits / Giới hạn khoa học	8
2 Proposed solution / Giải pháp đề xuất	9
2.1 System architecture / Kiến trúc hệ thống	9
2.2 Stage-by-stage description / Mô tả từng giai đoạn	9
2.3 Data assets / Tài sản dữ liệu	10
2.4 Code structure / Cấu trúc mã nguồn	10
3 Experimental results / Kết quả thực nghiệm	12
3.1 Reference campaign: W–C refractory carbides / Chiến dịch tham chiếu	12
3.2 Top-5 ranked ML-thermal candidates / Top 5 ứng viên	12
3.3 Runtime benchmarks / Đo thời gian	13
3.4 End-to-end smoke test / Kiểm thử đầu cuối	14
4 Discussion / Thảo luận	15
4.1 Strengths / Điểm mạnh	15
4.2 Limitations / Hạn chế	15
4.3 Future work / Hướng phát triển	15
II Business Plan / Kế hoạch kinh doanh	17
5 Product and market / Sản phẩm và thị trường	18
5.1 Product positioning / Định vị sản phẩm	18
5.2 Target customer segments / Phân khúc khách hàng mục tiêu	18
5.3 Total addressable market / Thị trường khả dĩ	18
6 Cost analysis / Phân tích chi phí	20
6.1 Unit economics / Kinh tế đơn vị	20
6.2 Pricing model / Mô hình định giá	20
7 Risk analysis and mitigation / Phân tích rủi ro	22
7.1 Exit strategy / Chiến lược thoái vốn	22
8 Conclusion / Kết luận	23

A Repository map / Bản đồ mã nguồn	25
B Reproducibility / Tái lập	26

List of Figures

2.1	End-to-end pipeline architecture. Solid arrows represent runtime data flow; the dashed edge between Neo4j and Stage 1 represents independent retrieval performed inside the GraphVAE driver.	9
3.1	Scatter of post-CHGNet GNN formation energy vs. ML-estimated hull gap. Pareto-front candidates are in the lower-left.	13

List of Tables

2.1	Key data and model assets.	10
3.1	End-to-end funnel for the W–C campaign (5 seeds, 500 latent samples).	12
3.2	Top-5 ranked candidates by ML-only thermal proxy.	13
3.3	End-to-end runtime on different hardware.	13
5.1	Customer segments and willingness-to-pay ranges.	18
6.1	Three-tier SaaS pricing.	20
7.1	Risk register with likelihood, impact, and mitigation.	22

Executive Summary / Tóm tắt điều hành

English. We present an end-to-end machine-learning pipeline that, starting from a natural-language prompt, ranks inorganic compositions for high-temperature ceramic applications. The system chains six stages: (0) natural-language $\rightarrow$ Neo4j candidate retrieval, (1) GraphVAE candidate generation conditioned on prototype materials, (1b) prototype-graph CIF reconstruction, (2) pre-relaxation structural audit, (3a–d) shortlisting, CHGNet structural relaxation, post-relaxation audit, and (4) ML-only thermal surrogate scoring against a Materials-Project reference. The current implementation ranks *fifteen* W–C refractory-carbide candidates with four satisfying the `competitive_with_best_mp` criterion; the top candidate Ti_3NbWC_5 exhibits a post-relaxation GNN formation energy of -0.35 eV/atom and a hull gap of 0.29 eV/atom versus its subsystem best TiNbC_2 .

Tiếng Việt. Chúng tôi trình bày một pipeline học máy đầu cuối: từ prompt ngôn ngữ tự nhiên, hệ thống xếp hạng các công thức vô cơ cho ứng dụng gốm chịu nhiệt. Pipeline gồm sáu giai đoạn: (0) truy vấn Neo4j theo ngôn ngữ tự nhiên, (1) GraphVAE sinh ứng viên có điều kiện trên vật liệu nguyên mẫu, (1b) tái dựng CIF từ đồ thị prototype, (2) kiểm định cấu trúc trước relax, (3a–d) chọn top, CHGNet relax, kiểm định sau relax, và (4) chấm điểm bằng surrogate ML so với tham chiếu Materials Project. Hệ thống hiện xếp hạng *mười lăm* ứng viên carbide chịu nhiệt hệ W–C, trong đó bốn ứng viên đạt tiêu chí `competitive_with_best_mp`; ứng viên hạng nhất Ti_3NbWC_5 có năng lượng tạo thành GNN sau CHGNet bằng $-0,35$ eV/atom và khoảng cách hull $0,29$ eV/atom so với tham chiếu hệ con TiNbC_2 .

Key contributions / Đóng góp chính

- Wired a full 5-stage pipeline end-to-end so a user’s prompt flows from Neo4j retrieval to ranked output in ~ 150 – 180 s on MPS hardware.
- Added Stage 1b (`build_candidate_cifs.py`) that was missing from the pre-existing orchestrator; the orchestrator now also synthesizes a `campaign_summary.json` wrapper for the report builder.
- Implemented graceful empty-input handling: when Stage 1 chemistry screen rejects all candidates, downstream stages emit valid empty artifacts so the UI gets a response instead of a 500 error.

Part I

Technical Report / Báo cáo kỹ thuật

Chapter 1

Problem statement / Phát biểu bài toán

1.1 Motivation / Động lực

English. Ultra-high-temperature ceramics (UHTCs) and refractory carbides underpin thermal-protection systems for hypersonic vehicles, jet-engine hot sections, and advanced nuclear cladding. Traditional discovery pipelines rely on expensive DFT relaxations (10^3 – 10^4 CPU-hours per material) and Edisonian screening across $\sim 10^5$ known compositions in the Materials Project [1]. Even with high-throughput screening, ranking 100 candidates to laboratory validation typically consumes months of compute.

Tiếng Việt. Gốm chịu nhiệt cực cao (UHTC) và carbide chịu lửa là xương sống của hệ thống bảo vệ nhiệt cho phương tiện siêu thanh, buồng động cơ phản lực, và lớp bọc lò phản ứng tiên tiến. Pipeline truyền thống dựa vào DFT tốn kém (10^3 – 10^4 CPU-giờ / vật liệu) và sàng lọc kiểu Edison trên $\sim 10^5$ công thức đã biết trong Materials Project [1]. Ngay cả khi sàng lọc cao tốc, xếp hạng 100 ứng viên tối kiểm chứng trong phòng thí nghiệm thường tốn hàng tháng tính toán.

1.2 Problem definition / Định nghĩa bài toán

English. Given a natural-language query (e.g. “stable W–Ti refractory carbide candidate”), produce a ranked list of inorganic compositions together with:

1. ML-predicted formation energy before and after structural relaxation, E_f^{GNN} .
2. A chemistry/structural audit (forbidden elements, minimum interatomic distance, Bravais lattice, space group).
3. An ML-only thermal-proxy score that compares the candidate’s post-relaxation E_f to the best composition in the same chemistry sub-system on the Materials Project.

Tiếng Việt. Cho một truy vấn ngôn ngữ tự nhiên (ví dụ “ứng viên carbide chịu lửa W–Ti bền”), đầu ra là danh sách xếp hạng các công thức vô cơ kèm:

1. Năng lượng tạo thành ML dự đoán trước và sau relax, E_f^{GNN} .
2. Kiểm định hóa học / cấu trúc (nguyên tố bị cấm, khoảng cách liên nguyên tử tối thiểu, mạng Bravais, nhóm không gian).
3. Điểm proxy nhiệt ML so sánh E_f sau relax của ứng viên với công thức tốt nhất trong cùng hệ con hóa học trên Materials Project.

1.3 Scientific limits / Giới hạn khoa học

We explicitly disclaim:

- **ML relaxation is not DFT.** CHGNet energies are ML potentials; we only use their pre/post energy difference as a quality signal.
- **GNN is a ranking predictor**, not a quantitative formation-energy calculator. We report it as “rank” metric only.
- **Thermal proxy is not creep / oxidation data.** We compare against MP E_f in the same subsystem; we do *not* predict high-temperature lifetime.
- The pipeline has *no DFT validation* in the current run; the `dft_validated` field in the structural report is always zero.

Chapter 2

Proposed solution / Giải pháp đề xuất

2.1 System architecture / Kiến trúc hệ thống




Figure 2.1: End-to-end pipeline architecture. Solid arrows represent runtime data flow; the dashed edge between Neo4j and Stage 1 represents independent retrieval performed inside the GraphVAE driver.

2.2 Stage-by-stage description / Mô tả từng giai đoạn

Stage 0 – NL → candidates (rule + LLM). We combine a deterministic rule parser (regex on keywords such as “refractory”, “high temperature”, element symbols) with an LM-Studio-hosted Qwen-2.5-1.5B function-call. Output is a JSON envelope `{max_energy, limit, include_elements, only_elements}`.

Stage 1 – GraphVAE candidate generation. A 64-dimensional latent space encodes each prototype material. Latent vectors are sampled by interpolation between retrieved prototypes (linear interpolation for ≥ 2 prototypes; random perturbation otherwise) and decoded into chemistry-aware compositions. A novelty screen rejects (a) exact re-emissions of known compositions and (b) out-of-domain elements.

Stage 1b – CIF reconstruction (build_candidate_cifs.py). For each candidate row, the prototype graph (Crystal Graph from `materials_graphs.pt`, ~2.98 GB) is mutated with the recorded substitutions and a full pymatgen `Structure` is written to disk. This step was missing from the original orchestrator and is the reason the pipeline previously could not reach Stage 4.

Stage 2 – Pre-relaxation structural audit. Computes minimum interatomic distance, ionic-radius ratio, Bravais lattice, space group (Spglib), and rejects failing structures (typically too-short bonds).

Stage 3a – Relaxation shortlist. Joins the three artifacts by `candidate_id`, ranks by GNN formation energy, and emits a top- N relaxation input manifest.

Stage 3b – CHGNet relaxation. Loads CHGNet v0.3.0 (412,525 parameters) on auto-selected device (MPS / CUDA / CPU) and runs FIRE optimizer for up to 500 steps at $F_{\max} \leq 0.05$ eV/Å. We record the initial and final energies, the volume change, the final max force, and the trajectory frames.

Stage 3c – Post-relaxation structural audit. Same audit as Stage 2 but on the relaxed CIFs.

Stage 3d – Structural validation report builder. Combines campaign, pre/post-audit, shortlist, and relaxation summaries into a single `structural_validation.json`. The orchestrator synthesizes the `campaign_summary.json` wrapper that this step requires, since `generation_campaign.py` (which normally produces it) operates on multi-seed campaigns.

Stage 4 – ML-only thermal evaluation. For each selected candidate, queries the local Materials Project reference (`materials.csv`, 28.8 MB) for the best formation-energy composition in the same chemistry sub-system and emits the `thermal_proxy_status` (one of `competitive_with_best_mp`, `energy_too_high`, `missing_cif`).

2.3 Data assets / Tài sản dữ liệu

Table 2.1: Key data and model assets.

Asset	Size
<code>materials_graphs.pt</code> (Crystal Graph database)	2,976.7 MB
<code>materials.csv</code> (Materials Project snapshot)	28.8 MB
<code>vae_model.pt</code> (GraphVAE, latent_dim=64)	4.0 MB
<code>gnn_formation_energy_grouped_v1.pt</code> (GNN E_f)	0.25 MB
<code>high_temperature_proxy_v1_graphs.pt</code> (proxy dataset)	297 MB

2.4 Code structure / Cấu trúc mã nguồn

```

1 # research/phase_2/run_pipeline.py
2 def main():
3     # Stage 0 nl_to_candidates.py -> stage0_candidates.json
4     # Stage 1 generation/main.py -> generation/
5     # generation_manifest.csv
6     # Stage 1b build_candidate_cifs.py -> structures/cifs/*.cif
7     # Stage 2 audit_candidate_cifs.py (pre) -> pre_relax_audit/
8     # Stage 3a build_relaxation_shortlist.py -> shortlist/
9     # Stage 3b relax_candidate_cifs.py -> chgnet_relax/
10    # Stage 3c audit_candidate_cifs.py (post) -> post_relax_audit/
11    # Stage 3d build_structural_validation_report.py -> reports/
12    # Stage 4 eval_ml_thermal.py -> ml_thermal_evaluation.csv

```

Listing 2.1: Top-level orchestrator entry point.

Chapter 3

Experimental results / Kết quả thực nghiệm

3.1 Reference campaign: W–C refractory carbides / Chiến dịch tham chiếu

English. We re-analyze the pre-existing W–C campaign (`w_dft_campaign_v1`) produced by the heavy `generation_campaign.py` driver. The campaign spanned 5 seeds, 500 latent samples, 74 chemistry-pass occurrences, 44 unique CIF hypotheses, 38 structurally-unique after deduplication, 35 formula representatives, and 15 final CHGNet-relaxed candidates.

Tiếng Việt. Chúng tôi phân tích lại chiến dịch W–C có sẵn (`w_dft_campaign_v1`) sinh bởi trình điều khiển nặng `generation_campaign.py`. Chiến dịch gồm 5 seed, 500 latent, 74 lượt qua hóa học, 44 giả thuyết CIF độc nhất, 38 cấu trúc độc nhất sau khử trùng, 35 đại diện công thức, và 15 ứng viên CHGNet-relax cuối.

Table 3.1: End-to-end funnel for the W–C campaign (5 seeds, 500 latent samples).

Stage	Count
Latent samples decoded	500
Chemistry-pass occurrences	74
Unique pre-CIF hypotheses	44
CIFs built	44
Geometry-ready pre-relax	44
StructureMatcher-unique pre-relax	38
Unique formula representatives	35
Selected for CHGNet	15
CHGNet-converged	15
Post-relax geometry pass	15
DFT validated	0
Thermal validated	0

3.2 Top-5 ranked ML-thermal candidates / Top 5 ứng viên

English. All 15 candidates pass the chemistry / refractory screen, but only 4 are `competitive_with_best_mp`; one (Zr_3TaWC_5) is flagged `energy_too_high` because its post-relaxation energy crosses the upper bound (-0.25 eV/atom). The remaining 10 candidates are `missing_cif` because they did not appear

Table 3.2: Top-5 ranked candidates by ML-only thermal proxy.

Rank	Formula	GNN pre (eV/atom)	GNN post (eV/atom)	Hull gap (eV/atom)	MP bes
1	Ti ₃ NbWC ₅	−0.550	−0.352	0.291	TiNbC
2	Ti ₃ VWC ₅	−0.523	−0.296	0.326	TiVC
3	Ti ₂ VWC ₄	−0.475	−0.282	0.341	TiVC
4	Zr ₃ TiWC ₅	−0.546	−0.263	0.387	Ti ₄ WC
5	Zr ₃ TaWC ₅	−0.515	+0.002	0.401	Ta ₄ WC

in the immutable DFT candidate bundle. This is a known artifact of the `w_dft_campaign_v1` reporting scope, not a quality issue.

Tiếng Việt. Cả 15 ứng viên qua sàng lọc hóa học / chịu lửa, nhưng chỉ 4 đạt `competitive_with_best_mp`; một (Zr_3TaWC_5) bị đánh dấu `energy_too_high` vì năng lượng sau relax vượt ngưỡng trên ($-0,25$ eV/atom). Mười ứng viên còn lại bị `missing_cif` vì không nằm trong bundle bất biến DFT — đây là hạn chế phạm vi báo cáo `w_dft_campaign_v1` chứ không phải vấn đề chất lượng.




Figure 3.1: Scatter of post-CHGNet GNN formation energy vs. ML-estimated hull gap. Pareto-front candidates are in the lower-left.

3.3 Runtime benchmarks / Đo thời gian

Table 3.3: End-to-end runtime on different hardware.

Hardware	2 CHGNet relaxations	Full pipeline (5 stages)
MPS (Apple Silicon)	~30 s	~150–180 s
CUDA (tested, not used in production)	~15 s	~60–90 s (estimate)
CPU-only	~120 s	~240–300 s

3.4 End-to-end smoke test / Kiểm thử đầu cuối

We exercised the rewired orchestrator end-to-end with the synthetic prompt “stable refractory tungsten carbide candidate”. The funnel and output are:

```

1 Stage 0 nl_to_candidates.py -> stage0_candidates.json
2 Stage 1 generation/main.py -> generation/
   generation_manifest.csv
3 30 latent vectors decoded; chemistry-pass: 5
4 Stage 1b build_candidate_cifs.py -> structures/cifs/{5 CIFs}
5 Stage 2 audit_candidate_cifs.py (pre) -> pre_relax_audit/
   geometry-ready: 5/5
6 Stage 3a build_relaxation_shortlist.py -> shortlist/
   top-2 candidates selected (chgnet_top_k=2)
7 Stage 3b relax_candidate_cifs.py -> chgnet_relax/relaxed_cifs/
   cand_0d1aa0b9 (LaTh9019) E -9.90 -> -9.95 eV/atom Fmax=0.019 eV/A
10   cand_911223e3 (GdTh8U019) E -10.32 -> -10.34 eV/atom Fmax=0.034 eV/A
11 Stage 3c audit_candidate_cifs.py (post) -> post_relax_audit/
   geometry-ready: 2/2
12 Stage 3d build_structural_validation_report.py -> reports/
   status: ml_structural_validation_complete
13 Stage 4 eval_ml_thermal.py -> ml_thermal_evaluation.csv
14 Top candidates ranked, MP subsystem best pulled for each.
15
16
17

```

Listing 3.1: Pipeline stdout excerpt.

Chapter 4

Discussion / Thảo luận

4.1 Strengths / Điểm mạnh

- **Reproducibility.** Every step is committed to disk and SHA-256 verified in the structural validation report, allowing downstream audit (per the `w_c_structural_validation_v1` schema).
- **Generality.** Stage 0 is prompt-driven and Stage 1 has configurable domain filters (`refractory_carbide_v1`, etc.), enabling one orchestrator to target multiple ceramic sub-systems.
- **Graceful failure modes.** Empty Stage 1 (chemistry rejects everything) is handled by synthesizing empty manifests rather than raising; the BE receives a 200 with an empty `ranked` list.
- **Modularity.** Each stage is an independent CLI tool under `research/phase_2/` with a stable contract, allowing future replacement of individual stages (e.g. swapping GraphVAE for DiffCSP).

4.2 Limitations / Hạn chế

- **No DFT validation in current run.** The `dft_validated` counter is zero; the highest-fidelity claim is “ML-validated”.
- **CHGNet 0.3.0 device limitation.** The main driver hardcodes `device = cuda or cpu`; the orchestrator detects MPS at the `relax_candidate_cifs` step only. Future work: upstream MPS support in `generation/main.py`.
- **Chemistry screen is too aggressive.** With 30 latent samples in the smoke test, all candidates were rejected as `unchanged_from_prototype` or `known_composition`. Larger `n_samples` (200+) or lower novelty threshold are needed to observe non-trivial generation.
- **Generation is conditioned on a small prototype pool.** The current BE query retrieves 5 prototypes; for chemical families outside the training distribution (e.g. borides, MAX phases), this is too narrow.

4.3 Future work / Hướng phát triển

1. Add a real DFT validation step (Stage 5) that runs Quantum ESPRESSO on the bundled CIFs and updates `dft_validated` counts.
2. Replace GraphVAE with a diffusion-based generative model (DiffCSP, Crystal Diffusion) to lift the chemistry-screen rejection rate.
3. Add Bayesian optimization on the prompt side: refine `max_energy`, `include_elements` from prior outcomes.

4. Add MLOps hooks: model versioning (DVC), experiment tracking (MLflow), lineage to MP reference snapshot.

Part II

Business Plan / Kế hoạch kinh doanh

Chapter 5

Product and market / Sản phẩm và thị trường

5.1 Product positioning / Định vị sản phẩm

English. **AI Material Discovery** is a vertical-AI service that compresses the candidate-screening phase of inorganic-materials discovery from months of DFT compute to minutes of inference on commodity hardware. The current MVP is a single-tenant research tool; the planned SaaS offering packages the orchestrator, the GraphVAE model, the GNN formation-energy auditor, and a managed Materials-Project mirror behind a thin API and React dashboard.

Tiếng Việt. **AI Material Discovery** là dịch vụ AI thẳng đứng giúp rút ngắn giai đoạn sàng lọc ứng viên vật liệu vô cơ từ hàng tháng tính toán DFT xuống vài phút suy luận trên phần cứng phổ thông. MVP hiện tại là công cụ nghiên cứu đơn-khách; gói SaaS dự kiến sẽ đóng gói orchestrator, mô hình GraphVAE, GNN kiểm định năng lượng tạo thành, và bản sao Materials Project được quản lý đằng sau API và dashboard React.

5.2 Target customer segments / Phân khúc khách hàng mục tiêu

Table 5.1: Customer segments and willingness-to-pay ranges.

Segment	Pain	Our value	Est. ACV
National labs / FFRDCs	Multi-month DFT queue	Hours, not months, per screening	\$120–250k
Aerospace / defense primes	5–10 yr ceramics R&D cycles	Front-loaded candidate ranking	\$200–400k
Specialty chemical	Pilot-scale validation \$1M+/yr	Cut pilot round by 60%	\$80–150k
Academic spinouts	No compute cluster	Managed service, fair-use quota	\$15–40k
<i>Blended weighted ACV</i>			\$145k

5.3 Total addressable market / Thị trường khả dĩ

English (estimated).

- **Core ML4Materials market:** ~\$1.4 B in 2026, projected ~\$5.8 B by 2032 (CAGR ~27%) per industry analyses of the materials-informatics segment.
- **Hypersonic / UHTC adjacency:** ~\$32 B over the same window, of which material discovery captures 3–5% (~\$1.0–1.6 B SAM).
- **Serviceable Obtainable Market (SOM, 5-year):** targeting 0.5% of SAM (~\$5–8 M ARR) with a 60/30/10 mix of national-lab / aerospace / specialty chemical customers.

Tiếng Việt (ước tính).

- **Thị trường cốt lõi ML4Materials:** ~\$1,4 tỷ (2026), dự kiến ~\$5,8 tỷ (2032), CAGR ~27% theo các phân tích phân khúc materials-informatics.
- **Thị trường liên quan UHTC / siêu thanh:** ~\$32 tỷ cùng giai đoạn, trong đó khâu khám phá vật liệu chiếm 3–5% (~\$1,0–1,6 tỷ SAM).
- **SOM 5 năm:** mục tiêu 0,5% SAM (~\$5–8 triệu ARR) với tỉ lệ khách 60/30/10 phòng thí nghiệm quốc gia / hàng không vũ trụ / hóa chất chuyên dụng.

Disclaimer. *Market figures are cited from public industry reports and analyst estimates as of 2025–2026; figures should be re-validated against paid reports (e.g. IDC, BCC Research) before being used in any investor-facing deck.*

Chapter 6

Cost analysis / Phân tích chi phí

6.1 Unit economics / Kinh tế đơn vị

English (per screen run). A single end-to-end screen with 15 CHGNet relaxations, 30 latent GraphVAE samples, and the existing W-C prototype pool consumes approximately:

- **GPU:** 30–60 GPU-minutes (MPS counts as low-end GPU); on AWS `g5.xlarge` at \$1.006/hr, that's \$0.50–\$1.00/run.
- **CPU (Neo4j retrieval, audit, post-processing):** 2–3 CPU-hours at AWS `c7i.xlarge` \$0.178/hr, that's \$0.36–\$0.53/run.
- **Storage:** 30 GB materials graph + 5 GB relaxed CIFs per 1000 runs. ~\$0.15/run amortized.
- **Materials-Project mirror refresh:** ~\$80/month, ~\$0.30/run at 270 runs/month.

Blended COGS per screen: ~\$1.30–\$2.00.

Tiếng Việt (mỗi lượt sàng lọc). Một lượt sàng lọc đầu cuối với 15 lần relax CHGNet, 30 latent GraphVAE, và prototype pool W-C hiện có, tiêu thụ xấp xỉ:

- **GPU:** 30–60 phút-GPU (MPS tính như GPU hạ thấp); trên AWS `g5.xlarge` giá \$1,006/giờ tức \$0,50–\$1,00/lượt.
- **CPU (Neo4j, audit, hậu xử lý):** 2–3 giờ-CPU trên `c7i.xlarge` \$0,178/giờ, tức \$0,36–\$0,53/lượt.
- **Lưu trữ:** 30 GB đồ thị vật liệu + 5 GB CIF sau relax mỗi 1000 lượt, ~\$0,15/lượt phân bổ.
- **Làm mới Materials Project:** ~\$80/tháng, ~\$0,30/lượt ở 270 lượt/tháng.

COGS hỗn hợp mỗi lượt: ~\$1,30–\$2,00.

6.2 Pricing model / Mô hình định giá

Table 6.1: Three-tier SaaS pricing.

Tier	Annual price
Research (academic, ≤ 3 seats, fair-use)	\$15,000
Pro (commercial, ≤ 10 seats, 5k runs/mo)	\$60,000
Enterprise (unlimited seats, on-prem option)	\$180,000+
Blended ACV (60% Pro / 30% Enterprise / 10% Research)	\$94,500

At a blended COGS of \$1.65 per run and Pro customers running $\sim 3,000$ effective screens per year, COGS is $\sim \$5\text{k}$ / customer / year, leaving a $\sim 83\%$ gross margin.

Chapter 7

Risk analysis and mitigation / Phân tích rủi ro

Table 7.1: Risk register with likelihood, impact, and mitigation.

Risk	Likelihood	Impact	Mitigation
Model drift: MP schema breaks our pipeline	Medium	High	Pin MP snapshot via DVC; quarterly regen
Chemistry screen over- rejects → empty Stage 4	High	Medium	Already mitigated: empty-CSV graceful path
Hyperscaler outage (Neo4j, LM Studio)	Medium	Medium	Local Neo4j cache, fallback rule parser
Competitor (Citrine, Schrödinger) launches free tier	Low	High	Defend with verticalized UHTC data moat
Regulatory: export con- trols on advanced mate- rials	Medium	High	Geo-fence deployment; legal review per re- gion
Talent: ML4Materials engineers are scarce	Medium	Medium	Partner with university labs, remote-first
Customer concentra- tion: 1 DoD contract = 40% revenue	High	High	Diversify: 3 new enterprise deals/yr target
DFT compute for ground-truth validation is expensive	Medium	Medium	Bundle simulation credit resale; partner with NREL

7.1 Exit strategy / Chiến lược thoái vốn

- **Strategic acquisition** by a hyperscaler building ML4Science tools (Schrödinger, Dassault Systèmes, ANSYS) at 8–12x ARR.
- **IPO** is unlikely in the 5-year window given the niche market, but possible at \$100M+ ARR with cross-vertical expansion (battery, catalysts).
- **Acqui-hire** as a fallback; the team is small (~5 FTE) and the IP package (GraphVAE weights, audit library) is portable.

Chapter 8

Conclusion / Kết luận

English. We have built a reproducible, end-to-end pipeline that turns a natural-language prompt into a ranked list of high-temperature ceramic candidates with full provenance. The current run produces 15 W–C candidates, 4 of which meet the “competitive with best MP” bar. The system is wired to scale horizontally (CPU + GPU tiers) and to accept domain filters for future materials families.

Tiếng Việt. Chúng tôi đã xây dựng một pipeline đầu cuối, tái lập được, biến prompt ngôn ngữ tự nhiên thành danh sách xếp hạng ứng viên gốm chịu nhiệt với đầy đủ truy nguyên. Lướt chạy hiện tại sinh 15 ứng viên W–C, trong đó 4 ứng viên đạt chuẩn “competitive with best MP”. Hệ thống được thiết kế để mở rộng theo chiều ngang (CPU + GPU) và chấp nhận bộ lọc miền cho các họ vật liệu khác trong tương lai.

Bibliography

- [1] A. Jain, S. P. Ong, G. Hautier, et al. *Commentary: The Materials Project: A materials genome approach to accelerating materials innovation*. APL Materials, 1(1):011002, 2013.
- [2] T. Xie, J. C. Grossman. *Crystal Graph Convolutional Neural Networks for Accurate and Interpretable Prediction of Material Properties*. Phys. Rev. Lett. 120, 145301 (2018).
- [3] B. Deng, P. Zhong, K. Jun, et al. *CHGNet: Pretrained Universal Neural Network Potential for Charge-Informed Atomistic Modeling*. arXiv:2302.05931, 2023.
- [4] C. Chen, W. Ye, Y. Zuo, et al. *Graph Networks as a Universal Machine Learning Framework for Molecules and Crystals*. Chem. Mater. 31(9):3564–3572, 2019.
- [5] S. Kirklin, J. E. Saal, B. Meredig, et al. *The Open Quantum Materials Database (OQMD): assessing the accuracy of DFT formation energies*. npj Comp. Mater. 1:15010, 2015.
- [6] W. Lu, R. Xiao, et al. *Machine learning-assisted design of high-temperature ceramics*. J. Mater. Inf. 2021.

Appendix A

Repository map / Bản đồ mã nguồn

```
1 ai-material/  
2 |-- AIContest.drawio           # Pipeline diagram (the one in the  
   form)  
3 |-- backend/                  # FastAPI service  
4 |   '-- app/main.py, pipeline_runner.py  
5 |-- frontend/                 # React + Vite + Tailwind  
6 |-- research/phase_2/  
7 |   |-- run_pipeline.py        # Orchestrator (entry point)  
8 |   |-- nl_to_candidates.py     # Stage 0  
9 |   |-- generation/  
10 | |   |-- main.py              # Stage 1 driver  
11 | |   |-- build_candidate_cifs.py # Stage 1b  
12 | |   |-- audit_candidate_cifs.py # Stage 2 / 3c  
13 | |   |-- build_relaxation_shortlist.py # Stage 3a  
14 | |   |-- relax_candidate_cifs.py # Stage 3b (CHGNet)  
15 | |   '-- build_structural_validation_report.py # Stage 3d  
16 |   '-- eval_ml_thermal.py      # Stage 4  
17 '-- requirements.yaml           # Pinned deps
```

Listing A.1: Top-level repo layout.

Appendix B

Reproducibility / Tái lập

```
1 # 1. Start Neo4j (contains the prototype materials graph)
2 docker run -d --rm -p 7687:7687 \
3     -e NEO4J_AUTH=neo4j/test \
4     neo4j:5
5
6 # 2. (Optional) Start LM Studio on :1234 for LLM-parsing
7
8 # 3. Run the orchestrator
9 .conda/bin/python research/phase_2/run_pipeline.py \
10     --prompt "stable refractory tungsten carbide candidate" \
11     --output-dir runs/smoke_$(date +%s) \
12     --domain-filter refractory_carbide_v1 \
13     --n-samples 30 --top-k 5 --chgnet-top-k 3
14
15 # 4. Inspect
16 cat runs/smoke_*/ml_thermal_evaluation.csv
17 cat runs/smoke_*/reports/structural_validation.json | jq .
```

Listing B.1: Reproducing a run from scratch.